

DLCV HW4 Report

Name: 黃宇平 Dep.:電信碩一 Student ID:R06942065

Problem 1: VAE

1. Describe the architecture & implementation details of your model(1%)


我的Encoder及Decoder架構分別如下:

Encoder:			
Layer (type)	Output Shape	Param #	Connected to
vae_in (InputLayer)	(None, 64, 64, 3)	0	
vae_enc_conv_0 (Conv2D)	(None, 32, 32, 64)	4864	vae_in[0][0]
batch_normalization_1 (BatchNor	(None, 32, 32, 64)	256	vae_enc_conv_0[0][0]
activation_1 (Activation)	(None, 32, 32, 64)	0	batch_normalization_1[0][0]
dropout_1 (Dropout)	(None, 32, 32, 64)	0	activation_1[0][0]
vae_enc_conv_1 (Conv2D)	(None, 16, 16, 128)	204928	dropout_1[0][0]
batch_normalization_2 (BatchNor	(None, 16, 16, 128)	512	vae_enc_conv_1[0][0]
activation_2 (Activation)	(None, 16, 16, 128)	0	batch_normalization_2[0][0]
dropout_2 (Dropout)	(None, 16, 16, 128)	0	activation_2[0][0]
vae_enc_conv_2 (Conv2D)	(None, 8, 8, 256)	819456	dropout_2[0][0]
batch_normalization_3 (BatchNor	(None, 8, 8, 256)	1024	vae_enc_conv_2[0][0]
activation_3 (Activation)	(None, 8, 8, 256)	0	batch_normalization_3[0][0]
dropout_3 (Dropout)	(None, 8, 8, 256)	0	activation_3[0][0]
flatten_1 (Flatten)	(None, 16384)	0	dropout_3[0][0]
vae_z_mean (Dense)	(None, 1024)	16778240	flatten_1[0][0]
vae_z_log_var (Dense)	(None, 1024)	16778240	flatten_1[0][0]
vae_z (Lambda)	(None, 1024)	0	vae_z_mean[0][0] vae_z_log_var[0][0]
Total params: 34,587,520 Trainable params: 34,586,624 Non-trainable params: 896			

Decoder:		
Layer (type)	Output Shape	Param #
vae_decoder_in (InputLayer)	(None, 1024)	0
dense_1 (Dense)	(None, 16384)	16793600
reshape_1 (Reshape)	(None, 8, 8, 256)	0
vae_dec_conv_2 (Conv2DTransp	(None, 16, 16, 256)	1638656
batch_normalization_4 (Batch	(None, 16, 16, 256)	1024
activation_4 (Activation)	(None, 16, 16, 256)	0
vae_dec_conv_1 (Conv2DTransp	(None, 32, 32, 128)	819328
batch_normalization_5 (Batch	(None, 32, 32, 128)	512
activation_5 (Activation)	(None, 32, 32, 128)	0
vae_dec_conv_0 (Conv2DTransp	(None, 64, 64, 64)	204864
batch_normalization_6 (Batch	(None, 64, 64, 64)	256
activation_6 (Activation)	(None, 64, 64, 64)	0
vae_decoded (Conv2D)	(None, 64, 64, 3)	4803
Total params: 19,463,043 Trainable params: 19,462,147 Non-trainable params: 896		

其中，兩者convolution layer後面皆有batch normalization，沒有pooling 或 up-sampling layers，另外encoder還有dropout。我的latent dimension是1024維，kl-loss-lambda是1e-5，影像讀入後先轉為介於[0, 1]間的值才輸入model。使用Early-stopping callback下39個Epoch便結束training。

2. Plot the learning curve (reconstruction loss & KL divergence) of your model(1%)




3. Plot 10 testing images and their reconstructed result of your model and report your testing MSE of the entire test set



(1%)

MSE of the entire test set: 0.0064

4. Plot 32 random generated images of your model(1%)



5. Visualize the latent space by mapping test images to 2D space (with tSNE) and color them with respect to an attribute of your choice(1%)

6. Discuss what you've observed and learned from implementing VAE(1%)

這次實作VAE之前，我先實作了Dense AE和Conv AE，效果為VAE >> Dense AE > Conv AE，其中VAE和Conv AE的差別只在於Encoder output及Decoder Input。可見'限制latent space distribution' 確實能為Auto-Encoder Performance帶來大幅度的提升。

Problem 2: GAN

1. Describe the architecture & implementation details of your model(1%)

我的Discriminator及Generator架構分別如下：

discriminator:		
Layer (type)	Output Shape	Param #
dcgan_discrim_in (InputLayer)	(None, 64, 64, 3)	0
dcgan_discrim_conv_0 (Conv2D)	(None, 32, 32, 128)	9728
leaky_re_lu_1 (LeakyReLU)	(None, 32, 32, 128)	0
dropout_1 (Dropout)	(None, 32, 32, 128)	0
dcgan_discrim_conv_1 (Conv2D)	(None, 16, 16, 256)	819456
leaky_re_lu_2 (LeakyReLU)	(None, 16, 16, 256)	0
dropout_2 (Dropout)	(None, 16, 16, 256)	0
dcgan_discrim_conv_2 (Conv2D)	(None, 8, 8, 512)	3277312
leaky_re_lu_3 (LeakyReLU)	(None, 8, 8, 512)	0
dropout_3 (Dropout)	(None, 8, 8, 512)	0
flatten_1 (Flatten)	(None, 32768)	0
dense_1 (Dense)	(None, 256)	8388864
leaky_re_lu_4 (LeakyReLU)	(None, 256)	0
dropout_4 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0
/usr/local/lib/python3.6/site-packages/keras/engine/training.py:975: UserWarning: Discrepancy between trainable weights and collected trainable weights, did you set `model.trainable` without calling `model.compile` after ? 'Discrepancy between trainable weights and collected trainable' Total params: 24,991,234 Trainable params: 12,495,617 Non-trainable params: 12,495,617		

generator:		
Layer (type)	Output Shape	Param #
dcgan_gen_in (InputLayer)	(None, 100)	0
dense_3 (Dense)	(None, 8192)	827392
batch_normalization_1 (Batch Normalization)	(None, 8192)	32768
activation_2 (Activation)	(None, 8192)	0
reshape_1 (Reshape)	(None, 8, 8, 128)	0
dcgan_gen_conv_0 (Conv2DTranspose)	(None, 16, 16, 512)	1638912
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 512)	2048
activation_3 (Activation)	(None, 16, 16, 512)	0
dcgan_gen_conv_1 (Conv2DTranspose)	(None, 32, 32, 256)	3277056
batch_normalization_3 (Batch Normalization)	(None, 32, 32, 256)	1024
activation_4 (Activation)	(None, 32, 32, 256)	0
dcgan_gen_conv_2 (Conv2DTranspose)	(None, 64, 64, 128)	819328
batch_normalization_4 (Batch Normalization)	(None, 64, 64, 128)	512
activation_5 (Activation)	(None, 64, 64, 128)	0
conv2d_1 (Conv2D)	(None, 64, 64, 3)	387
Total params: 6,599,427 Trainable params: 6,581,251 Non-trainable params: 18,176		

其中，兩者皆沒有pooling 或 up-sampling layers，generator有batch normalization 及dropout，Discriminator則無，但是最後有兩層分別為256和1的Dense layer。

我的latent dimension是100維，影像讀入後先轉為介於 $[-1, 1]$ 間的值才輸入discriminator。(Generator最後的activation為tanh)

2. Plot the learning curve (in the way you prefer) of your model and briefly explain what you think it represents(1%)
3. Plot 32 random generated images of your model(1%)



4. Discuss what you' ve observed and learned from implementing GAN(1%)

我發現在不調整loss metric的情形下，DCGAN很難Train得起來，且Model Structure稍有變動結果便會差很多。此外，一般DCGAN Train好後即便能成功生成圖片，所生成的圖仍會非常接近，這可能是因為Mode Collapse的關係。

5. Compare the difference between image generated by VAE and GAN, discuss what you' ve observed(1%)

VAE生成的結果不論是在圖像品質或是變化上都較一般的GAN好，且參數較少、收斂速度較快，因此VAE在Generation上並不失為一個好的Baseline。

Problem 3: ACGAN

1. Describe the architecture & implementation details of your model(1%)

我的Discriminator及Generator架構分別如下：


discriminator:			
Layer (type)	Output Shape	Param #	Connected to
acgan_discrim_in (InputLayer)	(None, 64, 64, 3)	0	
acgan_discrim_conv_0 (Conv2D)	(None, 32, 32, 128)	9728	acgan_discrim_in[0][0]
leaky_re_lu_1 (LeakyReLU)	(None, 32, 32, 128)	0	acgan_discrim_conv_0[0][0]
dropout_1 (Dropout)	(None, 32, 32, 128)	0	leaky_re_lu_1[0][0]
acgan_discrim_conv_1 (Conv2D)	(None, 16, 16, 256)	819456	dropout_1[0][0]
leaky_re_lu_2 (LeakyReLU)	(None, 16, 16, 256)	0	acgan_discrim_conv_1[0][0]
dropout_2 (Dropout)	(None, 16, 16, 256)	0	leaky_re_lu_2[0][0]
acgan_discrim_conv_2 (Conv2D)	(None, 8, 8, 512)	3277312	dropout_2[0][0]
leaky_re_lu_3 (LeakyReLU)	(None, 8, 8, 512)	0	acgan_discrim_conv_2[0][0]
dropout_3 (Dropout)	(None, 8, 8, 512)	0	leaky_re_lu_3[0][0]
flatten_1 (Flatten)	(None, 32768)	0	dropout_3[0][0]
dense_1 (Dense)	(None, 256)	8388864	flatten_1[0][0]
leaky_re_lu_4 (LeakyReLU)	(None, 256)	0	dense_1[0][0]
dropout_4 (Dropout)	(None, 256)	0	leaky_re_lu_4[0][0]
dense_2 (Dense)	(None, 1)	257	dropout_4[0][0]
dense_3 (Dense)	(None, 14)	3598	dropout_4[0][0]
activation_1 (Activation)	(None, 1)	0	dense_2[0][0]
activation_2 (Activation)	(None, 14)	0	dense_3[0][0]
/usr/local/lib/python3.6/site-packages/keras/engine/training.py:975: UserWarning: Discrepancy between weights and collected trainable weights, did you set 'model.trainable' without calling 'model.compile'? Discrepancy between trainable weights and collected trainable' Total params: 24,998,430 Trainable params: 12,499,215 Non-trainable params: 12,499,215			

generator:			
Layer (type)	Output Shape	Param #	Connected to
acgan_gen_class_in (InputLayer)	(None, 14)	0	
acgan_gen_in (InputLayer)	(None, 100)	0	
dense_4 (Dense)	(None, 100)	1500	acgan_gen_class_in[0][0]
multiply_1 (Multiply)	(None, 100)	0	acgan_gen_in[0][0] dense_4[0][0]
dense_5 (Dense)	(None, 8192)	827392	multiply_1[0][0]
batch_normalization_1 (BatchNormalizatio	(None, 8192)	32768	dense_5[0][0]
activation_3 (Activation)	(None, 8192)	0	batch_normalization_1[0][0]
reshape_1 (Reshape)	(None, 8, 8, 128)	0	activation_3[0][0]
acgan_gen_conv_0 (Conv2DTranspose)	(None, 16, 16, 512)	1638912	reshape_1[0][0]
batch_normalization_2 (BatchNormalizatio	(None, 16, 16, 512)	2048	acgan_gen_conv_0[0][0]
activation_4 (Activation)	(None, 16, 16, 512)	0	batch_normalization_2[0][0]
acgan_gen_conv_1 (Conv2DTranspose)	(None, 32, 32, 256)	3277056	activation_4[0][0]
batch_normalization_3 (BatchNormalizatio	(None, 32, 32, 256)	1024	acgan_gen_conv_1[0][0]
activation_5 (Activation)	(None, 32, 32, 256)	0	batch_normalization_3[0][0]
acgan_gen_conv_2 (Conv2DTranspose)	(None, 64, 64, 128)	819328	activation_5[0][0]
batch_normalization_4 (BatchNormalizatio	(None, 64, 64, 128)	512	acgan_gen_conv_2[0][0]
activation_6 (Activation)	(None, 64, 64, 128)	0	batch_normalization_4[0][0]
conv2d_1 (Conv2D)	(None, 64, 64, 3)	387	activation_6[0][0]
Total params: 6,600,927 Trainable params: 6,582,751 Non-trainable params: 18,176			

Convolutional Layers的維度、kernel數和DCGAN維持一致。其中，兩者皆沒有pooling 或 up-sampling layers，generator有batch normalization，Discriminator則有dropout。

Discriminator的output先經過一層Dense(256)，再分別通過Dense(14) 和 Dense(1) 得到Class prediction和 Real/Fake Label predictionGenerator的部分，Class label input先通過Dense(100) 再與Latent vector以點乘的方式Merge，之後便和DCGAN一樣Reshape後通入Convolution layers。

2. Plot the learning curve (in the way you prefer) of your model and briefly explain what you think it represents(1%)



在上圖中，我畫出每個 Epoch Discriminator和Generator 的Accuracy，由此圖大約可以看出兩個Accuracy始終維持互相對抗的關係，並沒有誰能上升至接近100%，且若一者上升，另一者有可能會略為下降。

3. Plot 10 pair of random generated images of your model, each pair generated from the same random vector input but with different attribute. This is to demonstrate your model's ability to disentangle feature of interest(2%)

